

Demeter - AI Engineer PRD

ML Ensemble & Simulation Engine

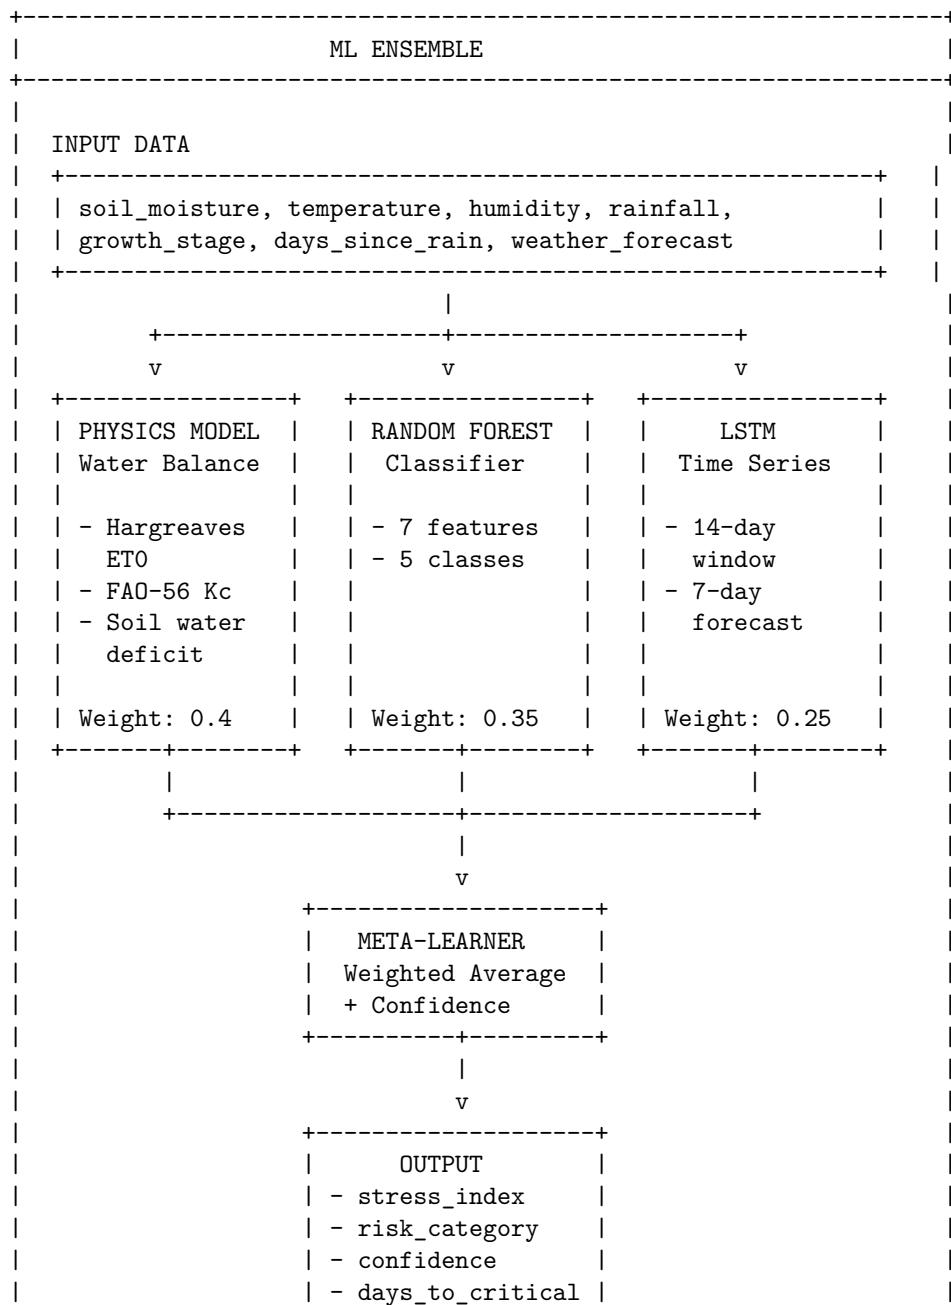
Role: AI/ML Engineer

Stack: Python 3.11 + FastAPI + scikit-learn + TensorFlow/Keras

1. Overview

Build the ML prediction service that combines physics-based models with machine learning to predict crop stress and simulate farming scenarios.

2. ML Ensemble Architecture



```

|           | - recommendation   |
|           | - forecast[]      |
|           +-----+
|
+-----+

```

3. API Endpoints

3.1 Prediction Endpoint

POST /predict

Request:

```
{
  "farm_id": "uuid",
  "current_data": {
    "soil_moisture": 32.5,
    "temperature": 34.0,
    "humidity": 45.0,
    "days_since_rain": 5
  },
  "farm_info": {
    "planting_date": "2026-02-01",
    "growth_stage": "VEGETATIVE",
    "latitude": 10.5105,
    "longitude": 7.4165
  },
  "weather_forecast": [
    {"date": "2026-02-26", "temp_max": 36, "temp_min": 22, "precipitation": 0},
    {"date": "2026-02-27", "temp_max": 37, "temp_min": 23, "precipitation": 0}
  ],
  "historical_data": [
    {"timestamp": "2026-02-24T10:00:00Z", "soil_moisture": 45, "temperature": 32},
    {"timestamp": "2026-02-25T10:00:00Z", "soil_moisture": 38, "temperature": 33}
  ]
}
```

Response:

```
{
  "stress_index": 72,
  "risk_category": "SEVERE",
  "confidence": 0.85,
  "days_to_critical": 4,
  "recommendation": "Irrigate 25mm within 3 days to prevent yield loss",
  "forecast": [
    {"day": 1, "stress": 72, "category": "SEVERE"},
    {"day": 2, "stress": 76, "category": "SEVERE"},
    {"day": 3, "stress": 79, "category": "SEVERE"},
    {"day": 4, "stress": 83, "category": "CRITICAL"},
    {"day": 5, "stress": 85, "category": "CRITICAL"},
    {"day": 6, "stress": 88, "category": "CRITICAL"},
    {"day": 7, "stress": 90, "category": "CRITICAL"}
  ],
  "model_contributions": {
    "physics": 0.4,
    "random_forest": 0.35,
  }
}
```

```
        "lstm": 0.25
    }
}
```

3.2 Simulation Endpoint

POST /simulate

Request:

```
{
  "farm_id": "uuid",
  "current_state": {
    "soil_moisture": 32.5,
    "stress_index": 45,
    "growth_stage": "VEGETATIVE",
    "days_since_planting": 30
  },
  "scenario": "DRY_WEEK",
  "parameters": {
    "duration_days": 14,
    "rainfall_mm": 0
  },
  "weather_forecast": [...]
}
```

Response:

```
{
  "baseline": {
    "stress_index": 45,
    "yield_impact": 0,
    "trajectory": [45, 46, 47, 48, 49, 50, 51]
  },
  "simulated": {
    "stress_index": 82,
    "yield_impact": -35,
    "trajectory": [45, 52, 58, 64, 70, 76, 82]
  },
  "recommendation": "Irrigate 25mm before day 3 to avoid critical stress",
  "risk_window": {
    "critical_day": 4,
    "action_deadline": 3
  }
}
```

3.3 Health Check

GET /health

4. Model Implementations

4.1 Physics Model: Water Balance

```
# models/water_balance.py

import numpy as np
from datetime import datetime, timedelta
```

```

class WaterBalanceModel:
    """
    FAO-56 based water balance model for crop stress estimation.
    """

    # Crop coefficients (Kc) by growth stage for maize
    CROP_COEFFICIENTS = {
        'EMERGENCE': 0.3,
        'VEGETATIVE': 0.7,
        'TASSELING': 1.2,
        'GRAIN_FILL': 1.0,
        'MATURITY': 0.6
    }

    # Soil parameters (sandy loam - common in Northern Nigeria)
    FIELD_CAPACITY = 0.28 # m³/m³
    WILTING_POINT = 0.10 # m³/m³
    ROOT_DEPTH = 0.6 # meters

    def calculate_et0(self, temp_max: float, temp_min: float,
                      latitude: float, day_of_year: int) -> float:
        """
        Hargreaves equation for reference evapotranspiration.
        ETO = 0.0023 x (T_mean + 17.8) x sqrt(T_max - T_min) x Ra
        """
        t_mean = (temp_max + temp_min) / 2
        t_range = max(temp_max - temp_min, 0.1)

        # Extraterrestrial radiation (simplified)
        ra = self._calculate_ra(latitude, day_of_year)

        et0 = 0.0023 * (t_mean + 17.8) * np.sqrt(t_range) * ra
        return max(et0, 0)

    def _calculate_ra(self, latitude: float, day_of_year: int) -> float:
        """
        Calculate extraterrestrial radiation (MJ/m²/day).
        """
        lat_rad = np.radians(latitude)
        solar_dec = 0.409 * np.sin(2 * np.pi * day_of_year / 365 - 1.39)
        ws = np.arccos(-np.tan(lat_rad) * np.tan(solar_dec))
        dr = 1 + 0.033 * np.cos(2 * np.pi * day_of_year / 365)

        ra = (24 * 60 / np.pi) * 0.082 * dr * (
            ws * np.sin(lat_rad) * np.sin(solar_dec) +
            np.cos(lat_rad) * np.cos(solar_dec) * np.sin(ws)
        )
        return ra

    def calculate_etc(self, et0: float, growth_stage: str) -> float:
        """
        Crop evapotranspiration: ETc = ETO x Kc
        """
        kc = self.CROP_COEFFICIENTS.get(growth_stage, 0.7)
        return et0 * kc

    def calculate_stress_factor(self, soil_moisture: float,
                               growth_stage: str) -> float:
        """
        """

```

```

Calculate stress factor Ks (0-1, where 1 = no stress).
"""
taw = (self.FIELD_CAPACITY - self.WILTING_POINT) * self.ROOT_DEPTH * 1000
raw = 0.5 * taw # Readily available water

current_water = (soil_moisture / 100) * self.ROOT_DEPTH * 1000
depletion = (self.FIELD_CAPACITY * self.ROOT_DEPTH * 1000) - current_water

if depletion <= raw:
    return 1.0 # No stress
elif depletion >= taw:
    return 0.0 # Max stress
else:
    return (taw - depletion) / (taw - raw)

def predict_stress_index(self, soil_moisture: float, temperature: float,
                        growth_stage: str, days_since_rain: int) -> float:
"""
Convert stress factor to 0-100 stress index.
"""
ks = self.calculate_stress_factor(soil_moisture, growth_stage)

# Base stress from water deficit
base_stress = (1 - ks) * 70

# Temperature stress (optimal 25-30degC for maize)
if temperature > 35:
    temp_stress = (temperature - 35) * 3
elif temperature < 15:
    temp_stress = (15 - temperature) * 2
else:
    temp_stress = 0

# Drought duration factor
duration_stress = min(days_since_rain * 2, 20)

# Growth stage sensitivity multiplier
stage_multipliers = {
    'EMERGENCE': 0.8,
    'VEGETATIVE': 0.9,
    'TASSELING': 1.3, # Most sensitive
    'GRAIN_FILL': 1.1,
    'MATURITY': 0.7
}
multiplier = stage_multipliers.get(growth_stage, 1.0)

stress_index = (base_stress + temp_stress + duration_stress) * multiplier
return min(max(stress_index, 0), 100)

def run_simulation(self, initial_state: dict, scenario: dict,
                  weather_forecast: list, days: int = 7) -> list:
"""
Simulate stress trajectory over time.
"""

trajectory = []
soil_moisture = initial_state['soil_moisture']

```

```

for day in range(days):
    weather = weather_forecast[day] if day < len(weather_forecast) else weather_forecast[-1]

    # Calculate water loss
    et0 = self.calculate_et0(
        weather['temp_max'],
        weather['temp_min'],
        initial_state.get('latitude', 10.0),
        datetime.now().timetuple().tm_yday + day
    )
    etc = self.calculate/etc(et0, initial_state['growth_stage'])

    # Apply scenario modifications
    rainfall = scenario.get('rainfall_mm', weather.get('precipitation', 0)) / days
    irrigation = scenario.get('irrigation_mm', 0) / days

    # Update soil moisture (simplified)
    soil_moisture += (rainfall + irrigation - etc * 0.5)
    soil_moisture = max(min(soil_moisture, 100), 5)

    # Calculate stress
    stress = self.predict_stress_index(
        soil_moisture,
        weather['temp_max'],
        initial_state['growth_stage'],
        initial_state.get('days_since_rain', 0) + day
    )

    trajectory.append({
        'day': day + 1,
        'stress': round(stress),
        'soil_moisture': round(soil_moisture, 1)
    })

return trajectory

```

4.2 Random Forest Classifier

```

# models/random_forest.py

import numpy as np
from sklearn.ensemble import RandomForestClassifier
import joblib
from pathlib import Path

class StressRandomForest:
    """
    Random Forest classifier for stress category prediction.
    """

    FEATURES = [
        'soil_moisture',
        'temperature',
        'humidity',
        'days_since_rain',
        'growth_stage_encoded',
        'cumulative_et',
    ]

```

```

        'forecast_rainfall_7d'
    ]

CATEGORIES = ['NONE', 'LOW', 'MODERATE', 'SEVERE', 'CRITICAL']

def __init__(self, model_path: str = None):
    if model_path and Path(model_path).exists():
        self.model = joblib.load(model_path)
    else:
        self.model = self._create_default_model()

def _create_default_model(self) -> RandomForestClassifier:
    """Create and train on synthetic data for hackathon."""
    model = RandomForestClassifier(
        n_estimators=100,
        max_depth=10,
        random_state=42
    )

    # Generate synthetic training data
    X, y = self._generate_synthetic_data(1000)
    model.fit(X, y)

    return model

def _generate_synthetic_data(self, n_samples: int):
    """Generate realistic synthetic training data."""
    np.random.seed(42)

    X = np.zeros((n_samples, len(self.FEATURES)))
    y = np.zeros(n_samples, dtype=int)

    for i in range(n_samples):
        soil_moisture = np.random.uniform(10, 80)
        temperature = np.random.uniform(20, 42)
        humidity = np.random.uniform(20, 90)
        days_since_rain = np.random.randint(0, 21)
        growth_stage = np.random.randint(0, 5)
        cumulative_et = np.random.uniform(20, 100)
        forecast_rain = np.random.uniform(0, 50)

        X[i] = [soil_moisture, temperature, humidity, days_since_rain,
                growth_stage, cumulative_et, forecast_rain]

    # Determine category based on features
    stress_score = (
        (80 - soil_moisture) * 0.4 +
        max(temperature - 30, 0) * 2 +
        days_since_rain * 1.5 -
        forecast_rain * 0.3
    )

    if stress_score < 15:
        y[i] = 0 # NONE
    elif stress_score < 30:
        y[i] = 1 # LOW
    elif stress_score < 50:

```

```

        y[i] = 2 # MODERATE
    elif stress_score < 70:
        y[i] = 3 # SEVERE
    else:
        y[i] = 4 # CRITICAL

    return X, y

def encode_growth_stage(self, stage: str) -> int:
    stages = ['EMERGENCE', 'VEGETATIVE', 'TASSELING', 'GRAIN_FILL', 'MATURITY']
    return stages.index(stage) if stage in stages else 1

def predict(self, features: dict) -> dict:
    """Predict stress category with probability."""
    X = np.array([
        features['soil_moisture'],
        features['temperature'],
        features['humidity'],
        features['days_since_rain'],
        self.encode_growth_stage(features['growth_stage']),
        features.get('cumulative_et', 50),
        features.get('forecast_rainfall_7d', 0)
    ])
    proba = self.model.predict_proba(X)[0]
    category_idx = np.argmax(proba)

    # Convert to stress index (center of category range)
    stress_ranges = [(0, 20), (21, 40), (41, 60), (61, 80), (81, 100)]
    stress_range = stress_ranges[category_idx]
    stress_index = (stress_range[0] + stress_range[1]) / 2

    return {
        'stress_index': stress_index,
        'category': self.CATEGORIES[category_idx],
        'confidence': float(proba[category_idx]),
        'probabilities': {cat: float(p) for cat, p in zip(self.CATEGORIES, proba)}
    }

def save(self, path: str):
    joblib.dump(self.model, path)

```

4.3 LSTM Time Series Model

```

# models/lstm.py

import numpy as np
from typing import List, Dict

try:
    import tensorflow as tf
    from tensorflow.keras.models import Sequential, load_model
    from tensorflow.keras.layers import LSTM, Dense, Dropout
    HAS_TF = True
except ImportError:
    HAS_TF = False

```

```

class StressLSTM:
    """
    LSTM model for time-series stress forecasting.
    """

SEQUENCE_LENGTH = 14 # 14 days of historical data
FORECAST_DAYS = 7 # Predict 7 days ahead

def __init__(self, model_path: str = None):
    if not HAS_TF:
        self.model = None
        return

    if model_path:
        try:
            self.model = load_model(model_path)
        except:
            self.model = self._create_model()
    else:
        self.model = self._create_model()

def _create_model(self):
    """Create LSTM architecture."""
    if not HAS_TF:
        return None

    model = Sequential([
        LSTM(64, input_shape=(self.SEQUENCE_LENGTH, 4), return_sequences=True),
        Dropout(0.2),
        LSTM(32, return_sequences=False),
        Dropout(0.2),
        Dense(16, activation='relu'),
        Dense(self.FORECAST_DAYS, activation='linear')
    ])

    model.compile(optimizer='adam', loss='mse', metrics=['mae'])

    # Quick train on synthetic data
    X, y = self._generate_sequences(500)
    model.fit(X, y, epochs=10, batch_size=32, verbose=0)

    return model

def _generate_sequences(self, n_samples: int):
    """Generate synthetic time-series data."""
    np.random.seed(42)

    X = np.zeros((n_samples, self.SEQUENCE_LENGTH, 4))
    y = np.zeros((n_samples, self.FORECAST_DAYS))

    for i in range(n_samples):
        # Generate a stress trajectory
        base_stress = np.random.uniform(20, 60)
        trend = np.random.uniform(-0.5, 1.5) # Increasing or decreasing
        noise = np.random.normal(0, 3, self.SEQUENCE_LENGTH + self.FORECAST_DAYS)

        full_trajectory = base_stress + trend * np.arange(self.SEQUENCE_LENGTH + self.FORECAST_DAYS) +

```

```

full_trajectory = np.clip(full_trajectory, 0, 100)

# Features: stress, soil_moisture (inverse), temperature, days
X[i, :, 0] = full_trajectory[:self.SEQUENCE_LENGTH] # stress
X[i, :, 1] = 80 - full_trajectory[:self.SEQUENCE_LENGTH] * 0.5 # soil moisture
X[i, :, 2] = 25 + full_trajectory[:self.SEQUENCE_LENGTH] * 0.1 # temperature
X[i, :, 3] = np.arange(self.SEQUENCE_LENGTH) # time

y[i] = full_trajectory[self.SEQUENCE_LENGTH:]

# Normalize
X = X / 100.0
y = y / 100.0

return X, y

def predict(self, historical_data: List[Dict]) -> Dict:
    """
    Predict future stress from historical data.

    historical_data: List of {timestamp, soil_moisture, temperature, stress_index}
    """
    if not HAS_TF or self.model is None:
        return self._fallback_predict(historical_data)

    # Prepare sequence
    sequence = np.zeros((1, self.SEQUENCE_LENGTH, 4))

    for i, data in enumerate(historical_data[-self.SEQUENCE_LENGTH:]):
        idx = i + max(0, self.SEQUENCE_LENGTH - len(historical_data))
        sequence[0, idx, 0] = data.get('stress_index', 50) / 100.0
        sequence[0, idx, 1] = data.get('soil_moisture', 50) / 100.0
        sequence[0, idx, 2] = data.get('temperature', 30) / 100.0
        sequence[0, idx, 3] = i / self.SEQUENCE_LENGTH

    # Predict
    forecast = self.model.predict(sequence, verbose=0)[0] * 100
    forecast = np.clip(forecast, 0, 100)

    return {
        'forecast': [{day: i+1, stress: round(float(s))} for i, s in enumerate(forecast)],
        'average_stress': float(np.mean(forecast)),
        'trend': 'increasing' if forecast[-1] > forecast[0] else 'decreasing'
    }

def _fallback_predict(self, historical_data: List[Dict]) -> Dict:
    """
    Simple fallback when TensorFlow not available.
    """
    if not historical_data:
        return {'forecast': [{'day': i+1, 'stress': 50} for i in range(7)],
                'average_stress': 50, 'trend': 'stable'}

    recent = [d.get('stress_index', 50) for d in historical_data[-7:]]
    avg = np.mean(recent)
    trend = (recent[-1] - recent[0]) / len(recent) if len(recent) > 1 else 0

    forecast = []
    for i in range(self.FORECAST_DAYS):

```

```

        stress = avg + trend * (i + 1)
        stress = max(0, min(100, stress))
        forecast.append({'day': i + 1, 'stress': round(stress)})

    return {
        'forecast': forecast,
        'average_stress': float(np.mean([f['stress'] for f in forecast])),
        'trend': 'increasing' if trend > 0.5 else 'decreasing' if trend < -0.5 else 'stable'
    }

def save(self, path: str):
    if self.model:
        self.model.save(path)

```

4.4 Ensemble Combiner

```

# models/ensemble.py

from typing import Dict, List
from .water_balance import WaterBalanceModel
from .random_forest import StressRandomForest
from .lstm import StressLSTM

class EnsemblePredictor:
    """
    Combines physics, ML, and deep learning models for robust predictions.
    """

    WEIGHTS = {
        'physics': 0.40,
        'random_forest': 0.35,
        'lstm': 0.25
    }

    def __init__(self):
        self.physics_model = WaterBalanceModel()
        self.rf_model = StressRandomForest()
        self.lstm_model = StressLSTM()

    def predict(self, request: Dict) -> Dict:
        """
        Generate ensemble prediction.
        """

        current = request['current_data']
        farm = request['farm_info']
        weather = request.get('weather_forecast', [])
        history = request.get('historical_data', [])

        # Physics model prediction
        physics_stress = self.physics_model.predict_stress_index(
            soil_moisture=current['soil_moisture'],
            temperature=current['temperature'],
            growth_stage=farm['growth_stage'],
            days_since_rain=current['days_since_rain']
        )

        # Random Forest prediction

```

```

rf_features = {
    'soil_moisture': current['soil_moisture'],
    'temperature': current['temperature'],
    'humidity': current['humidity'],
    'days_since_rain': current['days_since_rain'],
    'growth_stage': farm['growth_stage'],
    'cumulative_et': current.get('cumulative_et', 50),
    'forecast_rainfall_7d': sum(w.get('precipitation', 0) for w in weather[:7])
}
rf_result = self.rf_model.predict(rf_features)

# LSTM prediction
lstm_result = self.lstm_model.predict(history)
lstm_stress = lstm_result['average_stress']

# Weighted ensemble
ensemble_stress = (
    self.WEIGHTS['physics'] * physics_stress +
    self.WEIGHTS['random_forest'] * rf_result['stress_index'] +
    self.WEIGHTS['lstm'] * lstm_stress
)

# Calculate confidence from model agreement
stresses = [physics_stress, rf_result['stress_index'], lstm_stress]
std_dev = np.std(stresses)
confidence = max(0.5, 1 - (std_dev / 50)) # Higher agreement = higher confidence

# Determine risk category
risk_category = self._get_risk_category(ensemble_stress)

# Calculate days to critical
days_to_critical = self._estimate_days_to_critical(
    ensemble_stress,
    lstm_result['forecast'],
    weather
)

# Generate recommendation
recommendation = self._generate_recommendation(
    ensemble_stress,
    risk_category,
    days_to_critical,
    current,
    farm
)

# Build forecast
physics_forecast = self.physics_model.run_simulation(
    {**current, **farm},
    {'rainfall_mm': sum(w.get('precipitation', 0) for w in weather)},
    weather
)

# Blend forecasts
blended_forecast = []
for i in range(7):
    phys = physics_forecast[i]['stress'] if i < len(physics_forecast) else ensemble_stress

```

```

        lstm = lstm_result['forecast'][i]['stress'] if i < len(lstm_result['forecast']) else ensemble_stress
        blended = self.WEIGHTS['physics'] * phys + (1 - self.WEIGHTS['physics']) * lstm
        blended_forecast.append({
            'day': i + 1,
            'stress': round(blended),
            'category': self._get_risk_category(blended)
        })

    return {
        'stress_index': round(ensemble_stress),
        'risk_category': risk_category,
        'confidence': round(confidence, 2),
        'days_to_critical': days_to_critical,
        'recommendation': recommendation,
        'forecast': blended_forecast,
        'model_contributions': {
            'physics': round(physics_stress),
            'random_forest': round(rf_result['stress_index']),
            'lstm': round(lstm_stress)
        }
    }

def simulate(self, request: Dict) -> Dict:
    """
    Run what-if simulation.
    """
    current = request['current_state']
    scenario = request['scenario']
    params = request['parameters']
    weather = request.get('weather_forecast', [])

    # Baseline (no changes)
    baseline_trajectory = self.physics_model.run_simulation(
        current,
        {'rainfall_mm': sum(w.get('precipitation', 0) for w in weather)},
        weather,
        days=params.get('duration_days', 7)
    )

    # Simulated scenario
    scenario_params = self._build_scenario_params(scenario, params)
    simulated_trajectory = self.physics_model.run_simulation(
        current,
        scenario_params,
        weather,
        days=params.get('duration_days', 7)
    )

    baseline_final = baseline_trajectory[-1]['stress']
    simulated_final = simulated_trajectory[-1]['stress']

    # Estimate yield impact
    yield_impact = self._estimate_yield_impact(simulated_final, current['growth_stage'])
    baseline_yield = self._estimate_yield_impact(baseline_final, current['growth_stage'])

    return {
        'baseline': {

```

```

'stress_index': baseline_final,
'yield_impact': baseline_yield,
'trajectory': [t['stress'] for t in baseline_trajectory]
},
'simulated': {
    'stress_index': simulated_final,
    'yield_impact': yield_impact,
    'trajectory': [t['stress'] for t in simulated_trajectory]
},
'recommendation': self._generate_simulation_recommendation(
    baseline_final, simulated_final, scenario, params
),
'risk_window': {
    'critical_day': next(
        (i+1 for i, t in enumerate(simulated_trajectory) if t['stress'] >= 80),
        None
    ),
    'action_deadline': max(1, next(
        (i for i, t in enumerate(simulated_trajectory) if t['stress'] >= 80),
        len(simulated_trajectory)
    ) - 1)
}
}

def _get_risk_category(self, stress: float) -> str:
    if stress <= 20: return 'NONE'
    if stress <= 40: return 'LOW'
    if stress <= 60: return 'MODERATE'
    if stress <= 80: return 'SEVERE'
    return 'CRITICAL'

def _estimate_days_to_critical(self, current_stress: float,
                               forecast: List[Dict], weather: List[Dict]) -> int:
    for i, f in enumerate(forecast):
        if f['stress'] >= 80:
            return i + 1
    return None # Not expected to reach critical

def _generate_recommendation(self, stress: float, category: str,
                            days_critical: int, current: Dict, farm: Dict) -> str:
    if category == 'NONE':
        return "Crop is healthy. Continue monitoring."
    elif category == 'LOW':
        return "Mild stress detected. Monitor soil moisture closely."
    elif category == 'MODERATE':
        return f"Moderate stress. Consider irrigating 15-20mm within {days_critical or 5} days."
    elif category == 'SEVERE':
        irrigation = 25 if farm['growth_stage'] == 'TASSELING' else 20
        return f"Severe stress! Irrigate {irrigation}mm within {days_critical or 3} days to prevent yield loss."
    else:
        return "Critical stress! Irrigate immediately (30mm+) to minimize crop damage."

def _build_scenario_params(self, scenario: str, params: Dict) -> Dict:
    if scenario == 'DRY_WEEK':
        return {'rainfall_mm': 0, 'duration_days': params.get('duration_days', 14)}
    elif scenario == 'IRRIGATION_TEST':
        return {'irrigation_mm': params.get('irrigation_mm', 20)}

```

```

    elif scenario == 'DELAYED_PLANTING':
        return {'delay_days': params.get('delay_days', 14)}
    return params

def _estimate_yield_impact(self, stress: float, growth_stage: str) -> int:
    """Estimate % yield loss based on stress and growth stage."""
    stage_sensitivity = {
        'EMERGENCE': 0.5,
        'VEGETATIVE': 0.7,
        'TASSELING': 1.5, # Most sensitive
        'GRAIN_FILL': 1.2,
        'MATURITY': 0.4
    }

    if stress <= 40:
        return 0

    multiplier = stage_sensitivity.get(growth_stage, 1.0)
    base_impact = (stress - 40) * 0.6 * multiplier

    return -min(round(base_impact), 80)

def _generate_simulation_recommendation(self, baseline: float, simulated: float,
                                         scenario: str, params: Dict) -> str:
    diff = simulated - baseline

    if diff <= 5:
        return "Minimal impact expected from this scenario."

    if scenario == 'DRY_WEEK':
        days = params.get('duration_days', 14)
        return f"A {days}-day dry period would increase stress by {round(diff)} points. Irrigate 25mm b"
    elif scenario == 'IRRIGATION_TEST':
        amount = params.get('irrigation_mm', 20)
        return f"Adding {amount}mm irrigation would reduce stress by {round(-diff)} points."
    elif scenario == 'DELAYED_PLANTING':
        delay = params.get('delay_days', 14)
        return f"Delaying planting by {delay} days would increase peak stress by {round(diff)} points."
    return f"This scenario increases stress by {round(diff)} points. Take preventive action."
}

import numpy as np # Add at top of file if not present

```

5. FastAPI Application

```

# app/main.py

from fastapi import FastAPI, HTTPException
from fastapi.middleware.cors import CORSMiddleware
from pydantic import BaseModel
from typing import List, Dict, Optional
import uvicorn

from models.ensemble import EnsemblePredictor

```

```

app = FastAPI(
    title="Demeter ML Service",
    description="AI-powered crop stress prediction and simulation",
    version="1.0.0"
)

app.add_middleware(
    CORSMiddleware,
    allow_origins=[ "*" ],
    allow_methods=[ "*" ],
    allow_headers=[ "*" ],
)
# Initialize ensemble
predictor = EnsemblePredictor()

class CurrentData(BaseModel):
    soil_moisture: float
    temperature: float
    humidity: float
    days_since_rain: int

class FarmInfo(BaseModel):
    planting_date: str
    growth_stage: str
    latitude: float = 10.5
    longitude: float = 7.4

class WeatherDay(BaseModel):
    date: str
    temp_max: float
    temp_min: float
    precipitation: float = 0

class HistoricalData(BaseModel):
    timestamp: str
    soil_moisture: float
    temperature: float
    stress_index: Optional[float] = None

class PredictionRequest(BaseModel):
    farm_id: str
    current_data: CurrentData
    farm_info: FarmInfo
    weather_forecast: List[WeatherDay] = []
    historical_data: List[HistoricalData] = []

class SimulationRequest(BaseModel):
    farm_id: str
    current_state: Dict
    scenario: str
    parameters: Dict
    weather_forecast: List[WeatherDay] = []

@app.get("/health")

```

```

def health_check():
    return {"status": "healthy", "service": "demeter-ml"}


@app.post("/predict")
def predict(request: PredictionRequest):
    try:
        result = predictor.predict(request.dict())
        return result
    except Exception as e:
        raise HTTPException(status_code=500, detail=str(e))

@app.post("/simulate")
def simulate(request: SimulationRequest):
    try:
        result = predictor.simulate(request.dict())
        return result
    except Exception as e:
        raise HTTPException(status_code=500, detail=str(e))

if __name__ == "__main__":
    uvicorn.run(app, host="0.0.0.0", port=8000)

```

6. Project Structure

```

ml-service/
  ||| app/
  |||   ||| __init__.py
  |||   ||| main.py          # FastAPI application
  |||   ||| config.py        # Configuration
  ||| models/
  |||   ||| __init__.py
  |||   ||| water_balance.py # Physics model
  |||   ||| random_forest.py # RF classifier
  |||   ||| lstm.py          # Time series model
  |||   ||| ensemble.py      # Ensemble combiner
  ||| synthetic/
  |||   ||| __init__.py
  |||   ||| data_generator.py # Synthetic data for training
  ||| tests/
  |||   ||| test_water_balance.py
  |||   ||| test_ensemble.py
  |||   ||| test_api.py
  ||| trained_models/        # Saved model weights
  |||   ||| rf_model.pkl
  |||   ||| lstm_model.h5
  ||| requirements.txt
  ||| Dockerfile
  ||| README.md

```

7. Dependencies

```
# requirements.txt
fastapi==0.109.0
uvicorn==0.27.0
pydantic==2.5.0
numpy==1.26.0
pandas==2.1.0
scikit-learn==1.4.0
tensorflow==2.15.0 # Optional, fallback exists
joblib==1.3.0
python-multipart==0.0.6
httpx==0.26.0
```

8. Deliverables Checklist

- Water balance model with Hargreaves ET0
 - Random Forest stress classifier
 - LSTM time series forecaster (with fallback)
 - Ensemble combiner with weighted averaging
 - FastAPI endpoints (/predict, /simulate, /health)
 - Synthetic data generator for training
 - Confidence scoring from model agreement
 - Recommendation generation
 - Yield impact estimation
 - Unit tests for models
 - Docker setup
-

9. Testing the API

```
# Start server
uvicorn app.main:app --reload --port 8000

# Test prediction
curl -X POST http://localhost:8000/predict \
-H "Content-Type: application/json" \
-d '{
    "farm_id": "test-farm",
    "current_data": {
        "soil_moisture": 32,
        "temperature": 34,
        "humidity": 45,
        "days_since_rain": 5
    },
    "farm_info": {
        "planting_date": "2026-02-01",
        "growth_stage": "VEGETATIVE",
        "latitude": 10.51,
        "longitude": 7.42
    }
}'
```

Coordinate with: Backend (API integration), Hardware (sensor data format)